

# Calibration for probabilistic classification

Nick Normandin

# Overview

# The problem

some classifiers produce  
well-calibrated probabilities

- ▶ discriminant analysis
- ▶ logistic regression

others don't

- ▶ naive bayes
- ▶ SVMs
- ▶ anything with boosting
- ▶ tree methods
- ▶ sometimes neural networks

# First of all, who cares?

1. people with asymmetric misclassification costs
2. people who are going to use the scores in post-processing
3. people who want to compare model outputs on a fair basis

# Definitions: “classification”

in general, a classifier is a mapping function  $f$  such that

$$f : \vec{x} \mapsto c$$

where  $\vec{x} \in \mathbb{R}^P$ , but we're mostly interested in the intermediate step in where the function produces some membership score  $s_i$  for each instance  $\vec{x}_i$

# Definitions: “well-calibrated”

- ▶ for a model  $f$  and score  $s_i$  to be well-calibrated for class  $c_i$ , the empirical probability of a correct classification  $P(c_i | f(c_i | x_i) = s_i)$  must converge to  $f(c_i | x_i) = s_i$
- ▶ **example:** when  $s_i = 0.9$ , the probability of a correct classification should converge to  $P(c_i | s_i = 0.9) = 0.9$ . Otherwise, this isn't *really* a ‘probability.’

# Definitions: “calibration”

the calibration process is a separate mapping such that

$$g : s_i \mapsto P(c_i | s_i)$$

**it's really important to note that we're fitting another model on top of our model output, where your feature matrix is just the vector of probability scores  $\vec{s}$  and the target variable is the vector of true class labels  $\vec{y} \in \{0, 1\}$**

# Visualizing calibration



# Visualizing calibration

```
# train SVM w/ linear kernel on Pima Indian Diabetes data
m <- train(x = PimaIndiansDiabetes[,1:8],
           y = PimaIndiansDiabetes[,9], tuneLength = 1,
           method = 'svmLinear',
           trControl = trainControl(method = 'cv',
                                     savePredictions = T,
                                     classProbs = T))

pred <- m$pred[order(m$pred$rowIndex),]
result <- data.table(prob = pred$pos,
                     class = ifelse(pred$obs == 'pos', 1, 0))
```


# Visualizing calibration

plotting the model class score  $s_i$  vs the true label  $y_i$ . Is this a useful representation?



# Reliability plots

**(1)** Bin predictions by  $s_i$  (x-axis), **(2)** calculate  $p(c_i)$  by bin (y-axis)



# Common methods

# Platt scaling

pass  $s_i$  through the sigmoid


$$P(c_i | s_i) = \frac{1}{1 + \exp(As_i + B)}$$

where  $A$  and  $B$  are the solution to

$$\operatorname{argmax}_{A,B} - \sum_i y_i \log(p_i) + (1 - y_i) \log(1 - p_i)$$

# Platt scaling

applied to the Pima Indian Diabetes scores



# Isotonic regression

a strictly-nondecreasing piecewise linear function  $m$ , where


$$y_i = m(s_i) + \epsilon$$

fit such that

$$\hat{m} = \operatorname{argmin}_z \sum_i (y_i - z(s_i))^2$$

# Isotonic regression


linear and isotonic regression fit to random noise with drift





# Isotonic regression

applied to the Pima Indian Diabetes scores




# Notes on applying calibration

- ▶ it's really easy to overfit
  - ▶ calibration partition
  - ▶ cross-validation
- ▶ isotonic regression is generally more flexible (and can closely approximate sigmoid)
- ▶ best technique is dependent on the family of model used to generate  $s_i$

# Bootstrap aggregated isotonic regression

make it smoother by aggregating and averaging over 1000 resampled isotonic regression fits



## Extensions to $k > 2$

# Probabilistic classification as a simplex

- ▶ if we view the task of probabilistic classification as a vector-valued function, we can visualize the co-domain of this task as assigning the location of a prediction in a regular (unit) simplex,  $\Delta^{K-1}$
- ▶ why is this hard when  $K > 2$ ?

# Probabilistic classification as a simplex

 $\Delta^1$  $\Delta^2$ 

trivial with  $\Delta^1$  because we're only concerned with one unknown value and its complement. With  $\Delta^{K>2}$  the simplex becomes a triangle, tetrahedron, five-cell, etc.

# Multi-class probability estimation




Figure 1: classification problem with  $k = 4$

**Strategy:** decompose into separate binary classification problems

- ▶ one vs. all
- ▶ all pairs

# One vs. all




Figure 2: *one vs. all* reduces to  $k - 1$  calibrations



# All pairs




Figure 3: *all pairs* reduces to  $\binom{K}{2}$  calibrations


# Combining multi-class probability estimates

- ▶ least squares: minimize squared error loss w/ non-negativity
- ▶ coupling (only *all pairs*): minimize log loss w/ non-negativity
- ▶ normalization (only *one vs all*): divide by sum of probabilities estimates

## Experimental results

# effect of calibration set size

from Niculescu-Mizil and Caruana, 2005



## effect of multi-class combination method

from Zadrozny and Elkan, 2002

Method	MSE	Error Rate
NB Normalization	0.0326	0.1672
NB Least-Squares	0.0319	0.1672
NB Coupling	0.0304	0.1715
PAV NB Normalization	0.0241	0.1498
PAV NB Least-Squares	0.0260	0.1498
PAV NB Coupling	0.0260	0.1512
BNB Normalization	0.0163	0.0963
BNB Least-Squares	0.0164	0.0958
BNB Coupling	0.0160	0.1023
PAV BNB Normalization	0.0150	0.0946
PAV BNB Least-Squares	0.0150	0.0946
PAV BNB Coupling	0.0149	0.0935

# boosting causes calibration issues

from Niculescu-Mizil and Caruana, 2005

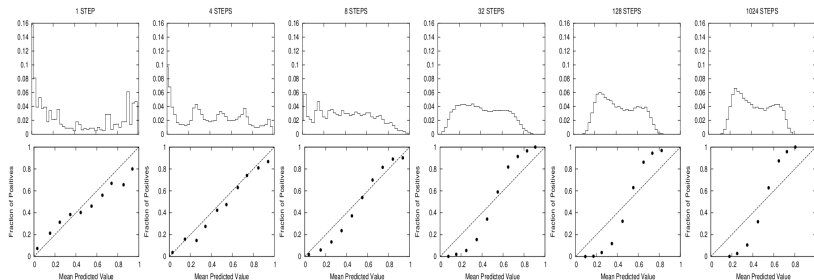


Figure 6: combination

# Conclusion

# References

- ▶ **Multivariate calibration of classifier scores into the probability space** by Martin Gebel
- ▶ **Transforming Classifier Scores into Accurate Multiclass Probability Estimates** by Zadrozny and Elkan
- ▶ **Obtaining Calibrated Probabilities from Boosting** by Niculescu-Mizil and Caruana
- ▶ **Predicting Good Probabilities With Supervised Learning** by Niculescu-Mizil and Caruana